

# CHAPTER 9



## *HashTables AND DICTIONARIES*

### LEARNING OBJECTIVES

- *Describe the purpose and use of a hashtable*
- *Describe the purpose and use of a dictionary*
- *Describe the purpose and use of key/value pairs*
- *Describe the purpose of a key*
- *List and describe the interfaces a class must implement if it's to be used as a key*
- *Describe the functionality provided by the IDictionary interface*
- *Use the Hashtable class in a program*
- *Use the Dictionary<TKey, TValue> class in a program*
- *Extract the key collection from a Hashtable or Dictionary*
- *Extract the value collection from a Hashtable or Dictionary*

---

## INTRODUCTION

---

Many times in your programming career you'll want to store an object in a collection and access that object via a unique *key*. This is the purpose of hashtables and dictionaries. Both hashtables and dictionaries perform the same type of service storing key/value pairs, but function differently on the insides. The `Hashtable` is the non-generic class and `Dictionary<TKey, TValue>` is its generic replacement.

In this chapter, I'm going to explain how hashtables work. Along the way I'll introduce you to concepts like *buckets*, *hash functions*, *keys*, *values*, *collisions*, and *growth factors*. I'll start the discussion with an overview of how hashtables work followed by a detailed demonstration of the operation of a chained hashtable with the help of a comprehensive coding example. In the end, you'll be able to create your own hashtable class, but that's the beauty of the .NET collections framework. You won't have to!

A few topics I'd like to cover in this chapter I'm going to put off until Chapter 10 — Coding For Collections. These include how to create immutable objects, how to write a custom comparer, and how to write a case insensitive hashcode provider, among others. You don't need to know these specialized topics just yet to fully use the `Hashtable` or `Dictionary<TKey, TValue>` classes. You will, however, need to read chapter 10 before you can use your own custom developed classes as hashtable or dictionary keys.

So, in the absence of violent objection, let's get started.

---

## HOW A HASHTABLE WORKS

---

A hashtable is an array that employs a special operation called a *hash function* to calculate an object's location within the array. The object being inserted into the table is referred to as the *value*. A hash function is applied to the value's associated *key* which results in an integer value that lies somewhere between the first array element (0) and the last array element (n-1). Figure 9-1 offers a simple illustration of a hashtable and a hash function being applied to a key.




Figure 9-1: Typical Hashtable Showing Hash Function Being Applied to a Key

Referring to figure 9-1 — the hashtable elements are referred to as *buckets*. The hash function transforms the key into an integer value that falls within the bounds of the array. Into this location the key's corresponding object is placed. All subsequent hashtable accesses using that particular key will "hash" to the same location.

## HASHTABLE COLLISIONS

A potential problem with hashtables arises when the hash function calculates the same hash value for two different keys. When this happens a *collision* is said to have occurred. There are several ways to resolve hashtable collisions and their complete treatment here is beyond the scope of this book, but if you're interested, I recommend you refer to Donald Knuth's excellent treatment of the subject. (Donald Knuth, *The Art of Computer Programming*, Vol. 3, Sorting and Searching, Second Edition)

I will, however, discuss and demonstrate one collision resolution strategy referred to as *chaining*. Chaining can be used to resolve hashtable collisions by storing values whose keys have hashed to the same bucket in a chain of elements. See figure 9-2.




Figure 9-2: Collision Resolution within a Chained Hashtable

Referring to figure 9-2 — collisions occur when two different keys hash to the same bucket location. The corresponding values are stored in a linked list or similar structure.

Collisions can be kept to a minimum if the hash function is especially good, meaning that for a large number of different keys the resulting hash values are uniformly distributed over the range of available bucket values. Also, the size of the hashtable can be increased as the number of keys grows to decrease the likelihood of a collision. These concepts are demonstrated in the following section.

## HomeGrownHashTable

The HomeGrownHashTable class presented in this section implements a complete example of a chained hashtable. The complete example consists of the KeyValuePair structure, presented in example 9.1, and the HomeGrownHashTable class, which is presented in example 9.2.

9.1 KeyValuePair.cs

```

1      using System;
2
3      public struct KeyValuePair : IComparable<KeyValuePair>, IComparable<string> {
4
5          private string _key;
6          private string _value;
7
8          public KeyValuePair(string key, string value){
9              _key = key;
10             _value = value;
11         }
12
13         public string Key {
14             get { return _key; }
15             set { _key = value; }
16         }
17
18         public string Value {
19             get { return _value; }
20             set { _value = value; }
21         }
22
23         public int CompareTo(KeyValuePair other){
24             return this._key.CompareTo(other.Key);
25         }
26
27         public int CompareTo(string other){
28             return this._key.CompareTo(other);
29         }
30     } // End KeyValuePair class definition

```

Referring to example 9.1 — the KeyValuePair structure represents the key/value pair that will be inserted into the hashtable. It defines two private fields named `_key` and `_value` and their corresponding public properties. The KeyValuePair structure also implements two interfaces: `IComparable<KeyValuePair>` and `IComparable<string>`. It implements these interfaces by defining two methods named `CompareTo(KeyValuePair other)` and `CompareTo(string)`.

other). The implementation of the `Comparable` interfaces is necessary if `KeyValuePair` objects are to be used in sorting and searching operations. For a more detailed discussion on how to prepare your custom structures and classes for use in collections see Chapter 10 — Coding for Collections.

Example 9.2 lists the code for the `HomeGrownHashtable` class. It uses the `KeyValuePair` structure as part of its implementation.

*9.2 HomeGrownHashtable.cs*

```

1      using System;
2      using System.Collections.Generic;
3      using System.IO;
4      using System.Text;
5
6      public class HomeGrownHashtable {
7
8          private float _loadFactor = 0;
9          private List<string> _keys = null;
10         private List<KeyValuePair>[] _table = null;
11         private int _tableSize = 0;
12         private int _loadLimit = 0;
13         private int _count = 0;
14
15         private const float MIN_LOAD_FACTOR = .65F;
16         private const float MAX_LOAD_FACTOR = 1.0F;
17
18         // Constructor methods
19         public HomeGrownHashtable(float loadFactor, int initialSize){
20             if((loadFactor < MIN_LOAD_FACTOR) || (loadFactor > MAX_LOAD_FACTOR)){
21                 Console.WriteLine("Load factor must be between {0} and {1}." +
22                     "Load factor adjusted to {1}", MIN_LOAD_FACTOR, MAX_LOAD_FACTOR);
23                 _loadFactor = MAX_LOAD_FACTOR;
24             } else {
25                 _loadFactor = loadFactor;
26             }
27
28             _keys = new List<string>();
29             _tableSize = this.DoublePrime(initialSize/2);
30             _table = new List<KeyValuePair>[_tableSize];
31             _loadLimit = (int)Math.Ceiling(_tableSize * _loadFactor);
32
33             for(int i = 0; i< _table.Length; i++){
34                 _table[i] = new List<KeyValuePair>();
35             }
36         } // end constructor
37
38
39         public HomeGrownHashtable():this(MAX_LOAD_FACTOR, 6){ }
40
41
42         public string[] Keys {
43             get { return _keys.ToArray(); }
44         } // end Keys property
45
46
47         public void Add(string key, string value){
48             if((key == null) || (value == null)){
49                 throw new ArgumentException("Key and Value cannot be null.");
50             }
51
52             string upperCaseKey = key.ToUpper();
53
54             if(_keys.Contains(upperCaseKey)){
55                 throw new ArgumentException("No duplicate keys allowed in HomeGrownHashtable!");
56             }
57
58             _keys.Add(upperCaseKey);
59             _keys.Sort();
60             int hashValue = this.GetHashValue(upperCaseKey);
61             KeyValuePair item = new KeyValuePair(upperCaseKey, value);
62             _table[hashValue].Add(item);
63             _table[hashValue].Sort();
64
65             if((++_count) >= _loadLimit){
66                 this.GrowTable();
67             }
68         } // end Add() method
69
70
71         public string Remove(string key){
72

```

```

73         if(key == null){
74             throw new ArgumentException("Key string cannot be null!");
75         }
76
77         string upperCaseKey = key.ToUpper();
78
79         if(_keys.Contains(upperCaseKey)){
80             _keys.Remove(upperCaseKey);
81         } else {
82             throw new ArgumentException("Key does not exist in table!");
83         }
84
85         _keys.Sort();
86         int hashValue = this.GetHashValue(upperCaseKey);
87         string return_value = string.Empty;
88         for(int i = 0; i < table[hashValue].Count; i++){
89             if(_table[hashValue][i].Key == upperCaseKey){
90                 return_value = _table[hashValue][i].Value;
91                 _table[hashValue].RemoveAt(i);
92                 _table[hashValue].Sort();
93                 break;
94             }
95         }
96
97         return return_value;
98     } // end Remove() method
99
100
101     private void GrowTable(){
102         List<KeyValuePair>[] temp = new List<KeyValuePair>[_table.Length];
103         for(int i=0; i<_table.Length; i++){
104             temp[i] = _table[i];
105         }
106
107         _table = new List<KeyValuePair>[this.DoublePrime(_tableSize)];
108
109         for(int i=0; i<temp.Length; i++){
110             _table[i] = temp[i];
111         }
112
113         for(int i=temp.Length; i<_table.Length; i++){
114             _table[i] = new List<KeyValuePair>();
115         }
116     } // end GrowTable() method
117
118
119     public string this[string key]{
120         get {
121             if((key == null) || (key == string.Empty)){
122                 throw new ArgumentException("Index key value cannot be null or empty!");
123             }
124             string return_value = string.Empty;
125             string upperCaseKey = key.ToUpper();
126             if(_keys.Contains(upperCaseKey)){
127                 int hashValue = this.GetHashValue(upperCaseKey);
128                 for(int i = 0; i < _table[hashValue].Count; i++){
129                     if(_table[hashValue][i].Key == upperCaseKey){
130                         return_value = _table[hashValue][i].Value;
131                         break;
132                     }
133                 }
134             }
135             return return_value;
136         }
137
138         set {
139             if((key == null) || (key == string.Empty)){
140                 throw new ArgumentException("Index key value cannot be null or empty!");
141             }
142
143             if((value == null) || (value == string.Empty)){
144                 throw new ArgumentException("String value cannot be null or empty!");
145             }
146             string upperCaseKey = key.ToUpper();
147             if(_keys.Contains(upperCaseKey)){
148                 int hashValue = this.GetHashValue(upperCaseKey);
149                 for(int i = 0; i < _table[hashValue].Count; i++){
150                     if(_table[hashValue][i].Key == upperCaseKey){
151                         KeyValuePair kvp = new KeyValuePair(upperCaseKey, value);
152                         _table[hashValue].RemoveAt(i);
153                         _table[hashValue].Add(kvp);

```

```

154         _table[hashValue].Sort();
155         break;
156     }
157 }
158 }
159 }
160 } // end indexer
161
162
163 private int DoublePrime(int currentPrime){
164     currentPrime *= 2;
165     int limit = 0;
166     bool prime = false;
167     while(!prime){
168         currentPrime++;
169         prime = true;
170         limit = (int)Math.Sqrt(currentPrime);
171         for(int i = 2; i<=limit; i++){
172             if((currentPrime % i) == 0){
173                 prime = false;
174                 break;
175             }
176         }
177     }
178     return currentPrime;
179 } // end DoublePrime() method
180
181
182 private int GetHashValue(string key){
183     int hashValue = ( Math.Abs(key.GetHashCode()) % _tableSize);
184     return hashValue;
185 } // end GetHashValue() method
186
187
188 public void DumpContentsToScreen(){
189     foreach(List<KeyValuePair> element in _table){
190         foreach(KeyValuePair kvp in element){
191             Console.Write(kvp.Value + " ");
192         }
193         Console.WriteLine();
194     }
195 } // end DumpContentsToScreen() method
196 } // end class definition

```

Referring to example 9.2 — the `HomeGrownHashtable` contains a `List<string>` object named `_keys` into which incoming key values are insert for future reference, and for the main table, a `List<KeyValuePair>` array named `_table`. The other fields include `_tableSize`, `_loadFactor`, `_loadLimit`, and `_count`. The *load factor* is used to calculate the *load limit*. In `HomeGrownHashtable`, the load factor is allowed to range between .65 and 1. When items are inserted into the hashtable, the calculated load limit is compared to the item count and if necessary, the main table is expanded to hold more elements.

The `HomeGrownHashtable` class defines the following methods: two constructors — one that does all the heavy lifting and a default constructor; `Add()`, `Remove()`, `GrowTable()`, `GetHashValue()`, `DoublePrime()`, and `DumpContentsToScreen()`. It also defines a `Keys` property and an indexer which allows values to be retrieved via their keys using familiar array notation. (i.e. `Hashtable_Reference["key"]`)

Let's step through the operation of the `Add()` method. The `Add()` method takes two arguments: a key string and a value string. The incoming arguments are checked for null values and if either are null the method throws an `ArgumentException`. The incoming key is converted to upper case with the `String.ToUpper()` method. The method then searches the `_keys` list to see if the key has already been inserted into the hashtable. If so, no duplicate keys are allowed and the method throws an `ArgumentException`. If the key is not in the `_keys` list, it's added to the list and the `_keys` list is then sorted. The key is then used to generate a hash value with the help of the `GetHashValue()` method. A new `KeyValuePair` object is created and added to the list at the `_table[hashValue]` location. That list is then sorted. The `_count` field is incremented and if necessary, the `_table` is expanded to hold additional elements by a call to the `GrowTable()` method.

Let's now examine the `GrowTable()` method. As its name implies, the purpose of the `GrowTable()` method is to grow the main hashtable (`_table`) to accommodate additional elements. The table growth mechanism is triggered in the `Add()` method when the element count (`_count`) approaches the hashtable's calculated load limit. The load limit (`_loadLimit`) is calculated in the body of the constructor: `_loadLimit = (int)Math.Ceiling(_tableSize * _loadFactor);` The first order of business in the `GrowTable()` method is to create a temporary `List<KeyValuePair>` array named `temp` and copy all the existing elements from `_table` to `temp`. A new array of `List<KeyValuePair>` elements is created dou-

ble the size of the existing table rounded up to the nearest prime number. This is done with the help of the `DoublePrime()` method. The reason I did this was because in my initial version of this example I used a custom hash function which relied on the generation of prime numbers to calculate the hash value of the key. I left the `DoublePrime()` method in the code so you can experiment with different hash generation techniques, most of which rely on prime numbers. (Note: The `DoublePrime()` method replaces the usual approach of maintaining an array of precalculated prime numbers.)

The `GetHashCode()` method calculates a hash value based on the key. Since I'm using strings as keys, I decided to rely on the `GetHashCode()` method defined by the `String` class. This value is then modded (%) with `_tableSize` to yield a value between 0 and `_tableSize - 1`. You can experiment with different hash generation formulas by replacing `key.GetHashCode()` with a custom hash generation function.

The indexer, which starts on line 119, allows values stored within `HomeGrownHashtable` to be accessed and set with familiar array notation using the key. It consists of two parts: the get and set sections. The get section checks the key to ensure its not null or the empty string. The key is converted to upper case and its existence is checked in the `_keys` list. If it's in the list, a hash value is generated and used to find the value's location within the `_table`. The `KeyValuePair` list located at that location must then be searched to find the key. When the key is found, the corresponding value is used to set `return_value`.

The set section works similar to the get section except that when the key is located, that `KeyValuePair` is removed and a new one created and added to the `KeyValuePair` list at that table location. The list is then sorted.

Example 9.3 offers a `MainApp` class that demonstrates the use of `HomeGrownHashtable`.

*9.3 MainApp.cs (Demonstrating HomeGrownHashtable)*

```
1      using System;
2
3      public class MainApp {
4          public static void Main(string[] args){
5              HomeGrownHashtable ht = new HomeGrownHashtable();
6              ht.Add("Rick", "Photographer, writer, publisher, handsome cuss");
7              ht.Add("Coralie", "Gorgeous, smart, funny, gal pal");
8              ht.Add("Kyle", "Tall, giant of a man! And a recent college graduate!");
9              ht.Add("Tati", "Thai hot sauce!");
10             Console.WriteLine(ht["Tati"]);
11             Console.WriteLine(ht["Kyle"]);
12             ht["Tati"] = "And a great cook, too!";
13             ht.DumpContentsToScreen();
14             ht.Remove("Tati");
15             ht.DumpContentsToScreen();
16         } // end Main() method
17     }
```

Referring to example 9.3 — an instance of `HomeGrownHashtable` is created on line 5. Lines 6 through 9 add several key/value pairs to hashtable. Lines 10 and 11 demonstrate the use of the indexer to access the values associated with the keys “Tati” and “Kyle”. On line 12, the indexer is used to replace the value associated with the key “Tati” with a new value. On line 13 the `DumpContentsToScreen()` method is called followed by the removal of the item referred to by the key “Tati” from the hashtable. The `DumpContentsToScreen()` method is then called one last time. Figure 9-3 shows the results of running this program.

```
C:\Collection Book Projects\Chapter_9\HomeGrownHashtable>mainapp
Thai hot sauce!
Tall, giant of a man! And a recent college graduate!
And a great cook, too!

Gorgeous, smart, funny, gal pal
Tall, giant of a man! And a recent college graduate!
Photographer, writer, publisher, handsome cuss

Gorgeous, smart, funny, gal pal
Tall, giant of a man! And a recent college graduate!
Photographer, writer, publisher, handsome cuss
C:\Collection Book Projects\Chapter_9\HomeGrownHashtable>
```

Figure 9-3: Results of Running Example 9.3



## Quick Review

A hashtable is an array that employs a special operation called a *hash function* to calculate an object's location within the array. The object being inserted into the table is referred to as the *value*. A hash function is applied to the value's associated *key* which results in an integer value that lies somewhere between the first array element (0) and the last array element (n-1).

The HomeGrownHashtable class implements a chained hashtable where each hashtable *bucket* points to a List<KeyValuePair> object into which KeyValuePair objects are inserted. The hashtable's load limit determines when the hashtable should be grown to accommodate additional elements. The load limit is calculated by multiplying the table size by the load factor.

Hashtable *collisions* can occur when two different keys hash to the same hash value. The chained hash table resolves collisions by allowing collisions to occur and storing the KeyValuePair objects in a list at that location which must then be searched to find the key/value pair of interest.

---

## Hashtable Class

---

The Hashtable class, located in the System.Collections namespace, stores hashtable elements as DictionaryEntry objects. A DictionaryEntry object consists of Key and Value properties and methods inherited from the System.Object class.

Unlike my HomeGrownHashtable discussed in the previous section, the Hashtable class doesn't use chaining to resolve collisions. According to Microsoft's documentation it uses a technique called *double hashing*. Double hashing works like this: If a key hashes to a bucket value already occupied by another key, the hash function is altered slightly and the key is rehashed. If that bucket location is empty the value is stored there, if it's occupied, the key must again be rehashed until an empty location is found.

Figure 9-4 shows the UML class diagram for the Hashtable inheritance hierarchy.

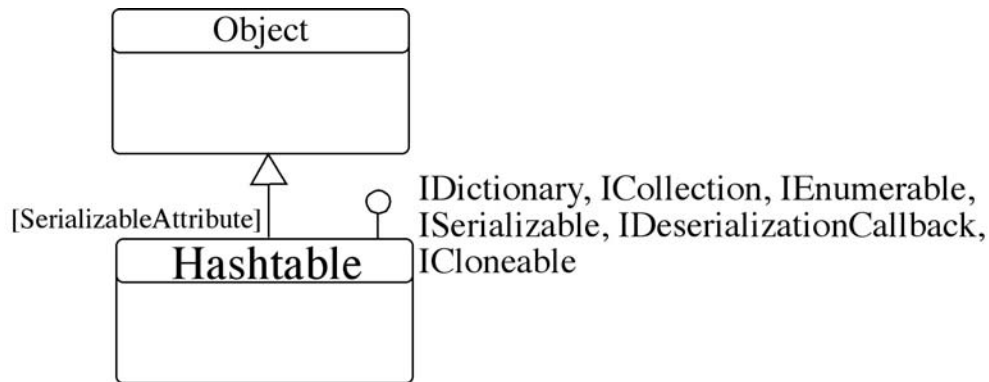


Figure 9-4: Hashtable Class Inheritance Hierarchy

Referring to figure 9-4 — The Hashtable class implements the IDictionary, ICollection, IEnumerable, ISerializable, IDeserializationCallback, and ICloneable interfaces. The functionality provided by each of these interfaces is discussed in more detail below.

### Functionality Provided by the IEnumerable Interface

The IEnumerable interface allows the items contained within the Hashtable collection to be iterated over with the foreach statement. Note that each element of a Hashtable collection is a DictionaryEntry object. The code to iterate over each element of a Hashtable with a foreach statement would look something like the following code snippet, assuming there exists a reference to a Hashtable object named ht:

```

foreach(DictionaryEntry item in ht){
    Console.WriteLine(item.Key);
}
  
```



The code snippet above would print out the value of each key to the console.

The Hashtable class also provides several properties, two of which are Keys and Values. The Keys property returns an ICollection of the Hashtable's keys and the Values property returns an ICollection of the Hashtable's Values. You can use a foreach statement to step through each of these collections as the following code snippet demonstrates, assuming that keys are strings:

```
foreach(string key in ht.Keys){
    Console.WriteLine(key);
}
```

In this example, each key in the Keys collection is written to the console. Later, in the Hashtable example program, I'll show you another way to step through the items in a Hashtable using its indexer.

Remember — when stepping through a collection with the foreach statement, you can't modify the elements.

### ***Functionality Provided by the ICollection Interface***

The ICollection interface declares the GetEnumerator() and CopyTo() methods as well as the IsSynchronized and SyncRoot properties. The IsSynchronized and SyncRoot properties are discussed in greater detail in Chapter 14 — Collections and Threads.

### ***Functionality Provided by the IDictionary Interface***

The IDictionary interface declares the Add(), Remove(), and Contains() methods, the indexer (shown as Item in the properties list), and the Keys and Values properties.

### ***Functionality Provided by the ISerializable and IDeserializationCallback Interfaces***

The ISerializable and IDeserializationCallback interfaces indicate that the Hashtable class requires custom serialization over and above simply tagging the class definition with the Serializable attribute. Custom serialization and deserialization is discussed in detail in Chapter 17 — Collections and I/O.

### ***Functionality Provided by the ICloneable Interface***

The ICloneable interface makes it possible to make copies of Hashtable objects.

## **Hashtable In Action**

Example 9.4 demonstrates the use of the Hashtable collection. In this example, the program reads a text file line-by-line and stores each line in the Hashtable collection using the line number as the key. The name of the text file must be supplied on the command line when the program is executed.

9.4 HashtableDemo.cs

```
1      using System;
2      using System.Collections;
3      using System.IO;
4      using System.Text;
5
6
7      public class HashtableDemo {
8
9          public static void Main(string[] args){
10              FileStream fs = null;
11              StreamReader reader = null;
12              Hashtable ht = new Hashtable();
13              try {
14                  fs = new FileStream(args[0], FileMode.Open);
15                  reader = new StreamReader(fs);
16
17                  int line_count = 1;
18                  string input_line = string.Empty;
19                  while((input_line = reader.ReadLine()) != null){
20                      string line_number_string = (line_count++).ToString();
21                      if(!ht.Contains(line_number_string)){
22                          ht.Add(line_number_string, input_line);
23                      }
24                  }
25              }
26              catch {
27              }
28          }
29      }
```

```

23     }
24 }
25
26 }catch(IndexOutOfRangeException){
27     Console.WriteLine("Please enter the name of a text file on the command line " +
28         "when running the program!");
29 }catch(Exception e){
30     Console.WriteLine(e);
31 } finally {
32     if(fs != null){
33         fs.Close();
34     }
35     if(reader != null){
36         reader.Close();
37     }
38 }
39
40
41 for(int i = 1; i<=ht.Keys.Count; i++){
42     Console.WriteLine("Line {0}: {1}", i, ht[i.ToString()]);
43 }
44
45 Console.WriteLine("*****");
46 Console.WriteLine("Line {0}: {1}", 2567, ht[2567.ToString()]);
47 Console.WriteLine("Line {0}: {1}", 193, ht[193.ToString()]);
48 Console.WriteLine("Line {0}: {1}", 669, ht[669.ToString()]);
49 Console.WriteLine("Line {0}: {1}", 733, ht[733.ToString()]);
50
51 } // end Main() method
52 } // end HashtableDemo class

```

Referring to example 9.4 — a `FileStream` reference named `fs` is declared on line 10 and initialized to null. On line 11 a `StreamReader` reference named `reader` is declared and also initialized to null. In the body of the try/catch block, which begins on line 13, the `FileStream` object is created using the first command line argument (`args[0]`). The `FileStream` object is then used to create the `StreamReader` object on the following line. On line 17, a local variable named `line_count` is declared and initialized to 1. An `input_line` variable of type string is declared on the following line and initialized to `string.Empty`. The while loop on line 19 processes each line of the text file. It first formulates a line number string (`line_number_string`) and checks its existence within the hashtable via the `Contains()` method. If it's not in the hashtable, the `input_line` is added using the `line_number_string` as its key.

The `for` statement on line 41 steps through each element of the hashtable using its indexer and writing the retrieved value to the console. Lines 46 through 49 access individual elements of the hashtable via the indexer.

The text file used for this example is named `Book.txt`. It contains the complete text of Cicero's *Tusculan Disputations*, by Marcus Tullius Cicero. It was downloaded from the Project Gutenberg website. ([www.gutenberg.net](http://www.gutenberg.net))

Figure 9-5 shows the results of running this program. Note that figure 9-5 only shows the last few lines of the output of 18517 lines of text printed to the console.

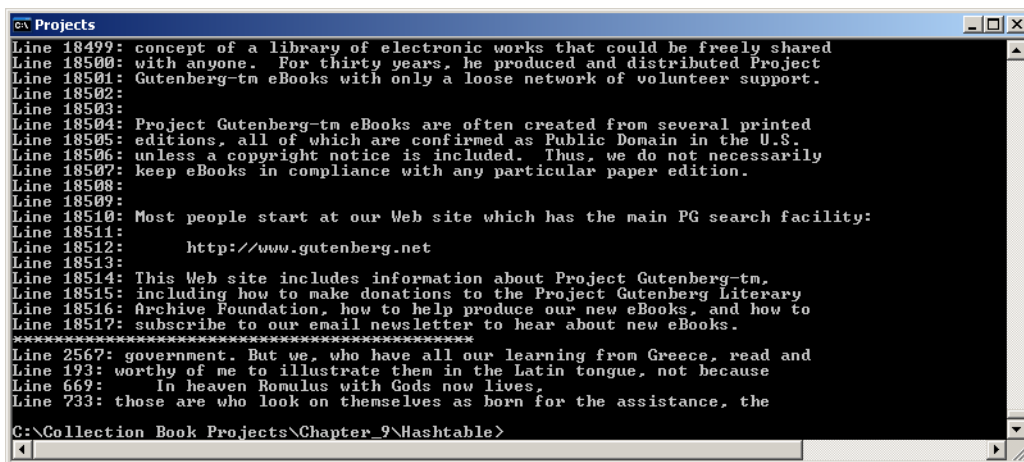


Figure 9-5: Results of Running Example 9.4

## Quick Review

The Hashtable is a non-generic collection class that stores its item as DictionaryEntry objects. The Hashtable class resolves collisions via *double hashing*, which is a process by which a key is rehashed using a modified hashing function until the collision has been resolved by the generation of a unique bucket location.

---

## Dictionary<TKey, TValue> Class

---

The Dictionary<TKey, TValue> class is the strongly-typed, generic version of the non-generic Hashtable class. The Dictionary<TKey, TValue> class also differs from the Hashtable class in the way it handles collisions. It uses chaining. Like the HomeGrownHashtable presented earlier, values whose keys hash to the same bucket are stored in a list, however, unlike the HomeGrownHashtable example, the Dictionary<TKey, TValue> class uses a different chain management algorithm which I'm positive is much more efficient than the approach I used.

Another huge difference between the Dictionary<TKey, TValue> class and the Hashtable class is the large number of extension methods provided by the System.Linq.Enumerable class that can be used on the Dictionary.

Figure 9-6 offers a UML class diagram showing the inheritance hierarchy of the Dictionary<TKey, TValue> class.

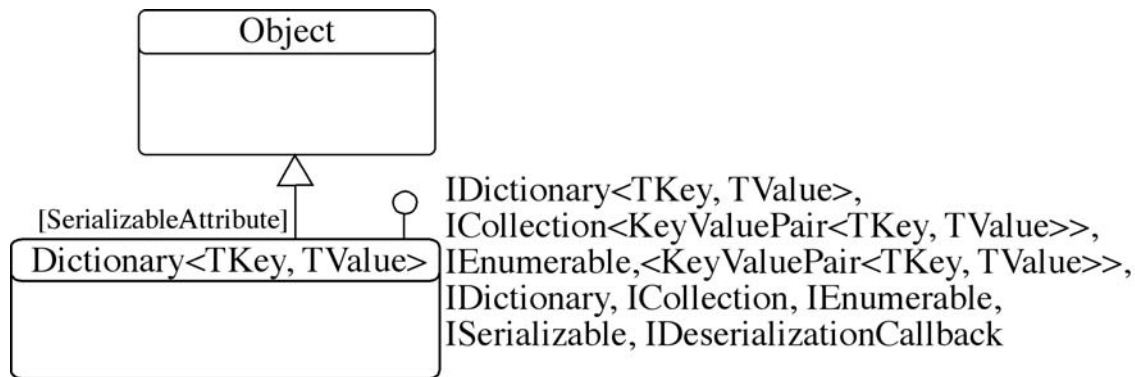


Figure 9-6: Dictionary<TKey, TValue> Class Inheritance Hierarchy

Referring to figure 9-6 — the Dictionary<TKey, TValue> class implements the IDictionary<TKey, TValue>, ICollection<KeyValuePair<TKey, TValue>>, IEnumerable<KeyValuePair<TKey, TValue>>, IDictionary, ICollection, IEnumerable, ISerializable, and IDeserializationCallback interfaces. Each of these interfaces are discussed in greater detail below.

### Functionality Provided by the IEnumerable and IEnumerable<KeyValuePair<TKey, TValue>> Interfaces

The IEnumerable and IEnumerable<KeyValuePair<TKey, TValue>> interfaces allow the items within a Dictionary<TKey, TValue> class to be iterated over with a `foreach` statement. Note that each element with a Dictionary<TKey, TValue> collection is a KeyValuePair<TKey, TValue> object. Assuming there was a reference to a Dictionary<string, int> object named `names_and_ages`, you could step through each element of the collection using a `foreach` statement similar to the following code snippet:

```
foreach (KeyValuePair<string, int> entry in names_and_ages) {
    Console.WriteLine("{0} is {1} years old!", entry.Key, entry.Value);
}
```

### Functionality Provided by the ICollection and ICollection<KeyValuePair<TKey, TValue>> Interfaces

The ICollection and ICollection<KeyValuePair<TKey, TValue>> interfaces tag the Dictionary<TKey, TValue> class as a collection type. The ICollection interface declares object synchronization properties `IsSynchronized` and `SyncRoot`, while the ICollection<KeyValuePair<TKey, TValue>> interface declares the `Add()`, `Remove()`, and `Con-`

tains() methods and the Count property. These interfaces also declare the GetEnumerator() methods required to iterate over the collection with a foreach statement.

### **Functionality Provided by the IDictionary and IDictionary<KeyValuePair<TKey, TValue>> Interfaces**

The IDictionary and IDictionary<KeyValuePair<TKey, TValue>> interfaces provide the non-generic and generic versions of Keys and Values properties, the indexer, and the ContainsKey() and the TryGetValue() methods.

### **Functionality Provided by the ISerializable and IDeserializationCallback Interfaces**

The ISerializable and IDeserializationCallback interfaces indicate that the Dictionary<TKey, TValue> collection requires custom serialization code over and beyond what the Serializable attribute alone provides.

### **Dictionary<TKey, TValue> Example**

Example 9.5 presents a short program demonstrating the use of the Dictionary<TKey, TValue> collection.

9.5 DictionaryDemo.cs

```

1      using System;
2      using System.Collections.Generic;
3      using System.Linq;
4
5      public class DictionaryDemo {
6
7          public static void Main(){
8              Dictionary<string, int> names_and_ages = new Dictionary<string, int>();
9              names_and_ages.Add("Rick", 49);
10             names_and_ages.Add("Kyle", 23);
11             names_and_ages.Add("Sport", 39);
12             names_and_ages.Add("Coralie", 39);
13             names_and_ages.Add("Tati", 21);
14             names_and_ages.Add("Schmoogle", 7);
15
16             foreach(KeyValuePair<string, int> entry in names_and_ages){
17                 Console.WriteLine("{0} is {1} years old!", entry.Key, entry.Value);
18             }
19
20             Console.WriteLine("The average age is {0:F4}", names_and_ages.Values.Average());
21
22         } // end Main() method
23     } // end DictionaryDemo class

```

Referring to example 9.5 — a Dictionary<string, int> reference named names\_and\_ages is declared and created on line 8. In this case, the keys will be strings and the values will be integers. Lines 9 through 14 add several entries into the dictionary. The foreach statement on line 16 steps through each KeyValuePair entry in the dictionary and prints the key and value to the console. Line 20 extracts the values from the dictionary via the Values property and calls the extension method Average() to calculate the average age. Figure 9-7 shows the results of running this program.




Figure 9-7: Results of Running Example 9.5

## Quick Review

The Dictionary<TKey, TValue> collection is a strongly-typed version of the Hashtable class that uses chaining instead of double hashing to resolve collisions. The Dictionary<TKey, TValue> collection stores its items as KeyValuePair objects. (See System.Collections.Generic.KeyValuePair<TKey, TValue> structure.)

---

## SUMMARY

---

A hashtable is an array that employs a special operation called a *hash function* to calculate an object's location within the array. The object being inserted into the table is referred to as the *value*. A hash function is applied to the value's associated *key* which results in an integer value that lies somewhere between the first array element (0) and the last array element (n-1).

The HomeGrownHashtable class implements a chained hashtable where each hashtable *bucket* points to a List<KeyValuePair> object into which KeyValuePair objects are inserted. The hashtable's load limit determines when the hashtable should be grown to accommodate additional elements. The load limit is calculated by multiplying the table size by the load factor.

Hashtable *collisions* can occur when two different keys hash to the same hash value. The chained hash table resolves collisions by allowing collisions to occur and storing the KeyValuePair objects in a list at that location which must then be searched to find the key/value pair of interest.

The Hashtable is a non-generic collection class that stores its item as DictionaryEntry objects. The Hashtable class resolves collisions via *double hashing*, which is a process by which a key is rehashed using a modified hashing function until the collision has been resolved by the generation of a unique bucket location.

The Dictionary<TKey, TValue> collection is a strongly-typed version of the Hashtable class that uses chaining instead of double hashing to resolve collisions. The Dictionary<TKey, TValue> collection stores its items as KeyValuePair objects. (See System.Collections.Generic.KeyValuePair<TKey, TValue> structure.)

---

## REFERENCES

---

Donald E. Knuth. The Art of Computer Programming, Vol. 1, Fundamental Algorithms. Third Edition. Addison-Wesley, Reading, Massachusetts. 1997. ISBN: 0-201-89683-4.

Microsoft Developer Network (MSDN) [<http://www.msdn.com>]

Arash Partow. The General Purpose Hash Function Algorithms Library (GeneralHashFunctionLibrary.java). [<http://www.partow.net/programming/hashfunctions/index.html>]

Project Gutenberg [<http://www.gutenberg.net>]

---

**NOTES**

---